

University of Amsterdam at CLEF 2026 Joker track

Jesse van Bakel, Robin Flier, Ezra Smink,
Jan Bakker and Jaap Kamps

University of Amsterdam

CLEF 2026



Task 1: Humor-aware IR

Approach

- Hybrid retrieval of BM25 + BGE-M3 raw combination
- Post processing with English RoBERTa pun classifier
- Reweighted with normalised F1 score
- RRF + mDeBERTa-v3 pun probability
- Evaluated with the MAP

Run	Description
1 UAms_Task1en_bm25	BM25 baseline (Anserini, stemming)
1 UAms_Task1en_bm25_filter40	BM25 + Filter on RoBERTa pun classifier (drops scores <40%)
1 UAms_Task1en_bm25_f1	BM25 + F1-score with RoBERTa Pun classifier score
1 UAms_Task1en_rm3	RM3 baseline (Anserini, stemming)
1 UAms_Task1en_rm3_filter40	RM3 + Filter on RoBERTa pun classifier (drops scores <40%)
1 UAms_Task1en_rm3_f1	RM3 + F1-score with RoBERTa Pun classifier score
1 UAms_Task1en_BGE_M3_Hybrid_00	BM25 + BGE-M3 (40/60) baseline
1 UAms_Task1en_BGE_M3_Hybrid_40	BM25 + BGE-M3 (40/60) + Filter on RoBERTa classifier
1 UAms_Task1en_BGE_M3_Hybrid_f1	BM25 + BGE-M3 (40/60) + F1-score with RoBERTa Pun classifier score
1 UAms_task1en_hybrid_norerank	BM25 + BGE-M3 (RRF), mDeBERTa-v3 pun detector
1 UAms_task1en_5000pool	BM25 + BGE-M3 (RRF) + BGE cross-encoder re-ranking
1 UAms_task1en_hybrid_norerank_wordnet0.3	BM25 + BGE-M3 (RRF) + mDeBERTa-v3 with WordNet polysemy
1 UAms_task1en_twostage_hw0.6_pool500	BM25 + BGE-M3 (RRF) for recall, mDeBERTa-v3 with WordNet ranking
1 UAms_Task1hi_bm25	BM25 baseline (Anserini, stemming)
1 UAms_Task1hi_rm3	RM3 baseline (Anserini, stemming)
1 UAms_Task1hi_BGE_M3_Hybrid_00	BGE-M3 baseline
1 UAms_Task1hi_BGE_M3_Hybrid_f1	BGE-M3 + F1-score with RoBERTa Pun classifier score

Task 1: Humor-aware IR

Results & Conclusion

- RRF with continuous pun probability scores best overall.
- Post processing with RoBERTa pun classifier and reweighted with normalised F1 score improves MAP over baseline

Run	MAP	GMAP	P@R	MRR	Precision				NDCG
					5	10	100	1,000	5
UAms_Task1en_bm25	10.87	0.00	7.85	20.80	6.99	4.72	2.82	0.37	10.35
UAms_Task1en_bm25_filter40	11.27	0.00	7.67	21.13	6.72	4.73	2.27	0.32	10.16
UAms_Task1en_bm25_f1	12.28	0.00	7.77	22.61	7.11	6.00	3.03	0.37	10.40
UAms_Task1en_rm3	10.35	0.00	6.78	18.69	6.37	4.81	2.85	0.38	9.13
UAms_Task1en_rm3_filter40	10.98	0.00	7.11	19.75	6.49	5.11	2.72	0.33	9.46
UAms_Task1en_rm3_f1	13.31	0.00	8.74	23.15	8.29	8.02	3.02	0.30	11.30
UAms_Task1en_BGE_M3_Hybrid_00	6.40	2.63	3.09	11.39	2.84	2.23	2.41	0.34	4.10
UAms_Task1en_BGE_M3_Hybrid_40	7.05	2.82	3.37	12.45	2.82	3.21	2.35	0.31	4.10
UAms_Task1en_BGE_M3_Hybrid_f1	8.75	3.74	5.60	15.88	5.11	6.00	2.46	0.34	4.10
UAms_task1en_hybrid_norerank	29.96	12.18	28.82	56.21	25.46	17.31	2.36	0.29	38.07
UAms_task1en_5000pool	21.91	9.29	20.91	37.67	19.96	16.64	2.36	0.29	26.07
UAms_task1en_hybrid_norerank_wordnet0.3	30.11	10.43	29.97	55.37	26.74	17.51	2.10	0.26	39.24
UAms_task1en_twostage_hw0.6_pool500	28.26	15.30	25.31	51.73	22.74	16.09	2.91	0.38	33.60
UAms_Task1hi_bm25	21.98	14.14	21.10	59.16	24.70	13.69	4.31	0.51	30.58
UAms_Task1hi_rm3	16.34	10.52	16.56	41.64	18.14	12.78	4.34	0.52	20.48
UAms_Task1hi_BGE_M3_Hybrid_00	5.27	3.26	2.71	12.61	2.97	2.02	3.53	0.48	3.93
UAms_Task1hi_BGE_M3_Hybrid_f1	4.05	2.93	0.76	4.94	0.78	0.74	3.63	0.48	3.93

Task 2: Wordplay Translation

Approach

- Finetuned translation models
- Multilingual vs French-specific classifier with human evaluation
- Reused mDeBERTa-v3-base pun detector from Task 1
- Compared larger translation backbones
- Evaluated with pun-location, BLEU and F1 score
- Extended to American Sign Language

2	UAms_task2_MarianMT_XLMRoBERTa	MarianMT + XLM-RoBERTa multilingual classifier
2	UAms_task2_MarianMT_CamemBERT	MarianMT + CamemBERT French-specific classifier
2	UAms_task2_marian_c0.0	MarianMT + mDeBERTa-v3 pun detector
2	UAms_task2_marian_c.25	MarianMT + content preservation
2	UAms_task2_nllb_c0.25	NLLB-1.3B + mDeBERTa-v3 pun detector and content preservation
2	UAms_task2_mbart_c.25	mBART + mDeBERTa-v3 pun detector and content preservation

Task 2: Wordplay Translation

Results & Conclusion

- Scaling up does not help
- Content already preserved
- CamemBERT scores better on pun classification
- Bottleneck: candidate pool
- Sign language output was not sufficiently interpretable

Model	Test Language	F1
XLM-RoBERTa	English	0.807
RoBERTa	English	0.794
XLM-RoBERTa	French	0.587
CamemBERT	French	0.694
XLM-RoBERTa	Spanish	0.394

#	team	# runs	Score	Rank	run_id
1	dsgt-arc	37	32.59	1	ensemble_0_0_0_100_gpt_pooled_transl
2	dsgt-arc	37	32.16	2	gpt_5.5_single
3	dsgt-arc	37	31.11	3	ensemble_0_0_0_100_all_pooled_transl
4	dsgt-arc	37	29.74	4	ensemble_0_0_0_100_claude_pooled_transl
5	dsgt-arc	37	28.88	5	ensemble_0_0_0_100_gemini_pooled_transl
6	n0shi	5	28.71	6	Satir_CVO_Direct_mistral-large:123b
7	arampageos	340	27.94	7	QANCHOR_TOP25
8	arampageos	340	27.90	8	NOTORCH_SELECTOR_STRICT
9	arampageos	340	27.90	9	QANCHOR_REFONLY
10	arampageos	340	27.90	10	NOTORCH_SELECTOR_MICRO5
11	arampageos	340	27.90	11	GROQ_FULL_PARTIAL_5
12	n0shi	5	25.76	325	Satir_Baseline
13	ezrasmink	5	25.22	336	UAms_marian-ft_content0.0
14	ezrasmink	5	25.22	337	UAms_marian-ft_content0.25
15	n0shi	5	25.11	342	Satir_BaselineCVO
16	n0shi	5	24.83	348	Satir_Classifier
17	robinflir	2	24.57	351	UAms_MarianMT_XLMRoBERTa
18	ezrasmink	5	24.47	355	UAms_opus-mt-en-fr-finetuned
19	robinflir	2	23.85	359	UAms_MarianMT_CamemBERT
20	ezrasmink	5	23.58	363	UAms_nllb-ft_content0.25
21	n0shi	5	22.70	368	Satir_ClassifierCVO
22	madhangi	1	22.29	370	PixelPair_groq_llama
23	ezrasmink	5	17.34	384	UAms_mbart-ft_content0.25
24	Baseline	1	7.05	391	identity_run

Task 3: Onomastic Wordplay Translation

Approach

- MarianMT translation model
- Rule based to scan the description on a cue
 - No cue = copy
 - No cue = MarianMT fallback
- Evaluated with exact match score only

3 Uams_task3_literal	MarianMT bare name (literal)
3 Uams_task3_context	MarianMT + description context
3 Uams_task3_copy_extract	Rule-based name extraction
3 Uams_task3_extended_context	Extraction + literal hybrid

Task 3: Onomastic Wordplay Translation

Results & Conclusion

- MarianMT without context performs worst
- Adding context improves slightly on the baseline
- Rarely a cue to fallback to

Table 1

Top-5 results per team for Task 3: Onomastic Wordplay Translation. **Score** is the official score; **Refs**, **Alter**, and **Copy** denote the percentages of reference matches, valid alternative translations, and untranslated outputs, respectively. **#runs** is the number of runs submitted by each team.

#	team	#runs	Score	Rank	Refs	Alter	Copy	run_id
1	UBO	5	71.88	1	44.11	45.37	8.20	gemini-3-flash
2	arampageos	41	43.13	2	51.58	7.35	40.38	V18_FRLEAD_STRICT_COPY
3	arampageos	41	43.13	3	51.58	7.35	40.38	V18_FRLEAD_WD_STRICT_COPY
4	arampageos	41	43.13	4	51.62	7.35	41.49	V14_OFFICIAL_B04_1700_2299
5	arampageos	41	43.13	5	51.58	7.35	40.38	V21_ASTERIX_DIRECT_ONLY
6	arampageos	41	43.13	6	51.58	7.35	40.38	V21_ASTERIX_COMBO
7	UBO	5	17.25	43	7.08	10.54	24.83	llama
8	ezrasmink	5	13.74	44	6.66	8.31	18.75	marian-ft_context
9	ezrasmink	5	13.74	45	6.97	8.31	18.55	cascade
10	Baseline	1	13.10	46	13.66	0.00	100.00	identity_run
11	ezrasmink	5	13.10	47	13.97	0.00	97.69	extraction
12	ezrasmink	5	7.35	49	5.47	2.24	21.21	marian-ft_barename
13	UBO	5	4.47	50	3.54	0.32	4.35	Qwen3-4B
14	pjmathematician	4	2.88	51	4.78	0.00	30.29	baseline_sub
15	UBO	5	2.24	55	2.42	1.60	6.31	mistral
16	UBO	5	1.92	56	1.35	0.64	0.15	gemma

Task 4: Humor Generation

Approach

- Instruction-tuned LLMs were asked to generate 6 candidates
- Reranked with:
 - The mDeBERTa-v3-base humour-detection classifier
 - Brief's content words present in the candidate
 - Ambiguity score
 - Distinctiveness score
- English generations evaluated by human annotators

Track	Run
EN	UAMS_task4_en_Qwen2.5-32B-instruct
EN	UAMS_task4_en_Qwen2.5-14B-instruct
EN	UAMS_task4_en_mistral-7B-instruct
EN	UAMS_task4_en_flan-t5-large
EN	UAMS_task4_en_flan-t5-base
ES	UAMS_task4_es-Llama-3-Instruct-Neur
ES	UAMS_task4_es_mistral-7B-instruct
ES	UAMS_task4_es_flan-t5-large
ES	UAMS_task4_es_flan-t5-base
ES	UAMS_task4_es-salamandra-7b-instruc
FR	UAMS_task4_fr-Chocolatine-2-14B
FR	UAMS_task4_fr_mistral-7B-instruct
FR	UAMS_task4_fr_flan-t5-large
FR	UAMS_task4_fr_flan-t5-base

Task 4: Humor Generation

Results & Conclusion


- Human evaluation: outputs are often natural sounding and correctly use the target pun word, but a lot less humorous
- Best model is Qwen2,5-32B-Instruct
- Larger models outperform smaller ones
- Encoder-decoder model underperform
- Reliable for structural correctness

Model	<i>n</i>	Natural English	Pun word used	Humorous
Qwen2.5-14B-Instruct	49	83.7 [71.0, 91.5]	100.0 [92.7, 100.0]	46.9 [33.7, 60.6]
Qwen2.5-32B-Instruct	49	73.5 [59.7, 83.8]	98.0 [89.3, 99.6]	55.1 [41.3, 68.1]
Mistral-7B-Instruct-v0.2	49	69.4 [55.5, 80.5]	95.9 [86.3, 98.9]	40.8 [28.2, 54.8]
flan-t5-large	50	84.0 [71.5, 91.7]	88.0 [76.2, 94.4]	10.0 [4.3, 21.4]
flan-t5-base	50	70.0 [56.2, 80.9]	76.0 [62.6, 85.7]	0.0 [0.0, 7.1]
All models	247	76.1	91.5	30.4

Table 8

CLEF 2026 JOKER Task 4: Official LLM Arena Results

	Rank	Run	BT	CI low	CI high	<i>n_b</i>
English	74	Qwen2.5-32B-Instruct	984	938	1025	316
	81	Qwen2.5-14B-Instruct	927	882	959	341
	84	Mistral-7B-Instruct-v0.2	922	882	952	341
	96	flan-t5-large	549	469	607	326
	98	flan-t5-base	422	328	466	327
Spanish	60	Llama-3-Instruct-Neurona-8b	1062	1027	1105	315
	113	Mistral-7B-Instruct-v0.2	903	858	936	339
	126	flan-t5-large	393	276	479	303
	127	flan-t5-base	344	272	406	340
	134	salamandra-7b-instruct	-270	-1734	-142	335
French	41	Chocolatine-2-14B-Instruct-v2.0.3	1191	1149	1230	332
	72	Mistral-7B-Instruct-v0.2	961	910	1001	314
	126	flan-t5-large	558	490	604	329
	128	flan-t5-base	508	442	567	330

The background is a solid medium purple color. It is decorated with several abstract elements: a cluster of white dots in the top-left corner, a large, irregular, semi-transparent white shape in the top-center, a smaller white dot cluster in the top-right, a white dot cluster in the bottom-left, and a large, solid, light-purple organic shape on the right side.

Thank you